

Project Title: Second Brain AI - System Architecture

1. Architectural Vision This project implements a **Local-First RAG (Retrieval Augmented Generation)** architecture. Unlike traditional cloud-native systems, this design prioritizes data privacy and zero-latency ingestion by performing vectorization on the edge (the user's machine) rather than relying on external API rate limits.

2. Core Components

- **Frontend:** React (Vite) + TailwindCSS for a responsive, real-time chat interface.
- **Backend:** FastAPI (Python) serving as the lightweight orchestration layer.
- **Vector Store: ChromaDB.** Chosen for its embedded nature, allowing the database to live within the application process. This removes the overhead of managing a separate database container (like Postgres) for a single-user application.
- **Embedding Model: all-MiniLM-L6-v2 (HuggingFace).** A highly efficient 384-dimensional model running locally on the CPU. This ensures user documents are never sent to third-party APIs for indexing, guaranteeing privacy.
- **LLM: Google Gemini Flash.** Used strictly for the final answer synthesis, ensuring high speed and low cost.

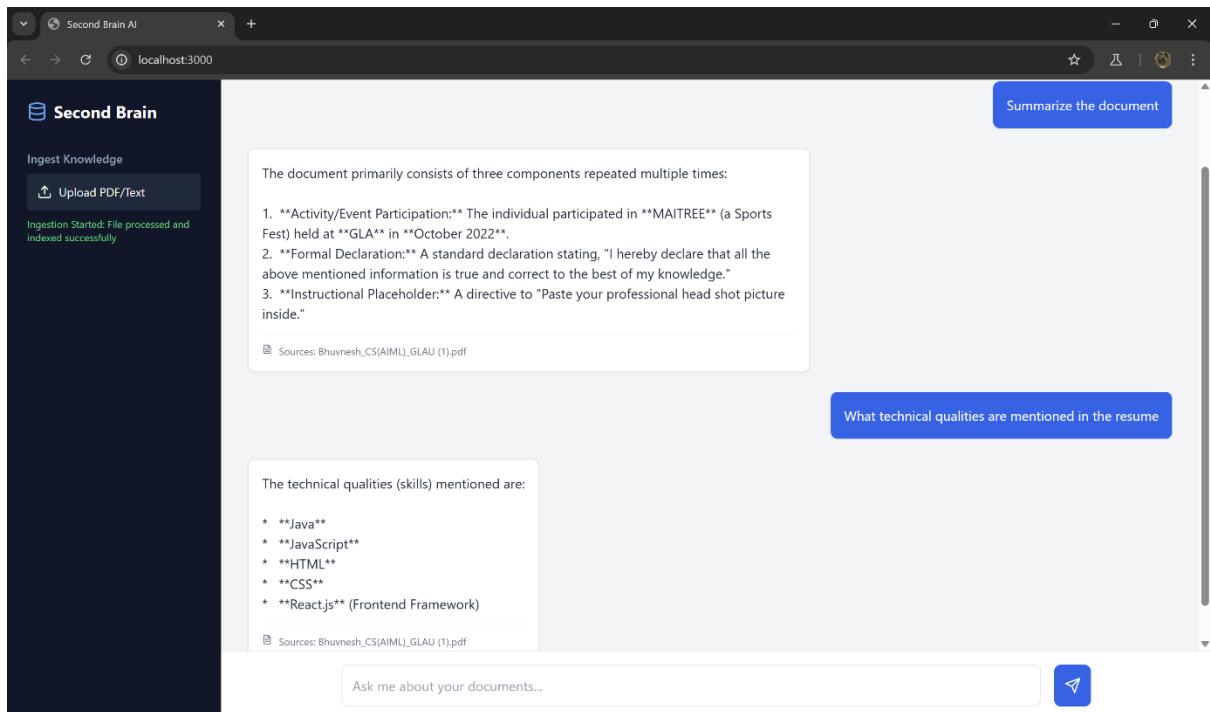
3. Data Flow

1. **Ingestion:** User uploads a PDF. The backend parses text using PyPDFLoader.
2. **Chunking:** Text is split into 500-character semantic chunks with 50-character overlap to preserve context.
3. **Local Indexing:** The local CPU converts chunks into vectors and stores them immediately in the persistent chroma_db on disk.
4. **Retrieval:** When a user asks a question, the system performs a Cosine Similarity search in ChromaDB to find the top 5 relevant chunks.
5. **Synthesis:** These chunks are injected into the Gemini system prompt to generate a grounded response.

4. Scalability & Trade-offs

- **Trade-off:** Processing is synchronous. For a massive multi-user system, I would introduce an asynchronous worker queue (Celery/Redis).
- **Decision:** For this "Personal Second Brain" MVP, the synchronous local model provides the fastest "Time-to-First-Token" and simplest deployment footprint.

Project Interface:



Terminal Log:

A screenshot of a terminal window from a code editor. The title bar shows tabs for PROBLEMS, OUTPUT, DEBUG CONSOLE, TERMINAL, and PORTS. The terminal tab is active, displaying the command "PS C:\Users\yash\Desktop\second_brain\backend> python main.py" and its output. The output shows the application starting a server, processing a PDF file, and handling several HTTP requests. The right side of the window shows a sidebar with icons for "python...", "esbuild fr...", and other development tools.

```
PS C:\Users\yash\Desktop\second_brain\backend> python main.py
vector_db = Chroma(
INFO:  Started server process [24544]
INFO:  Waiting for application startup.
INFO:  Application startup complete.
INFO:  Uvicorn running on http://0.0.0.0:8000 (Press CTRL+C to quit)
Processing Bhuvnesh_CS(AIML)_GLAU (1).pdf...
Success! Indexed 8 chunks.
INFO:  127.0.0.1:62300 - "POST /upload HTTP/1.1" 200 OK
INFO:  127.0.0.1:62310 - "OPTIONS /chat HTTP/1.1" 200 OK
INFO:  127.0.0.1:62310 - "POST /chat HTTP/1.1" 200 OK
INFO:  127.0.0.1:62312 - "POST /chat HTTP/1.1" 200 OK
```